




# parallel computing



speedup,  $S = T_1 / T_p$ , time  
with  $p$  processors vs with one

efficiency,  $E = T_1 / p T_p$

scalable algorithm –  $E$   
increasing function of  $n/p$   
where  $n$  is ‘problem size’



# parallel programming paradigms

shared memory

— partition work

$F(w_p)$ :

shared  $a$

lock( $a[i]$ )

work( $w_p$ )

unlock( $a[i]$ )

message passing

— partition data

$F(p)$ :


$a_p = a[p \dots p+(n/p)-1]$

work( $w$ )

exchange data( $a_p$ )

shared + partition data; message-passing + partition work also possible  
map-reduce: message-passing, data-parallel, pipelined work, higher level

# map-reduce



mappers:

take in  $k_1, v_1$  pairs

emit  $k_2, v_2$  pairs

$k_2, v_2 \leftarrow \text{map}(k_1, v_1)$

reducers:

receive all pairs for some  $k_2$

combine these in some manner


$k_2, f_r(\dots v_2 \dots) \leftarrow \text{reduce}(k_2, [\dots v_2 \dots])$

map-reduce platform responsible for *routing* pairs to reducers  
map-reduce reads data and *writes* fresh data; is a *batch* process

# map-reduce

Map:  $(d_k, 'w_1 \dots w'_n) \rightarrow [(w_i, c_i)]$  document -> word-count pairs

Reduce:  $(w_i, [c_i]) \rightarrow (w_i, \sum_i c_i)$  word, count-list -> word-count-total

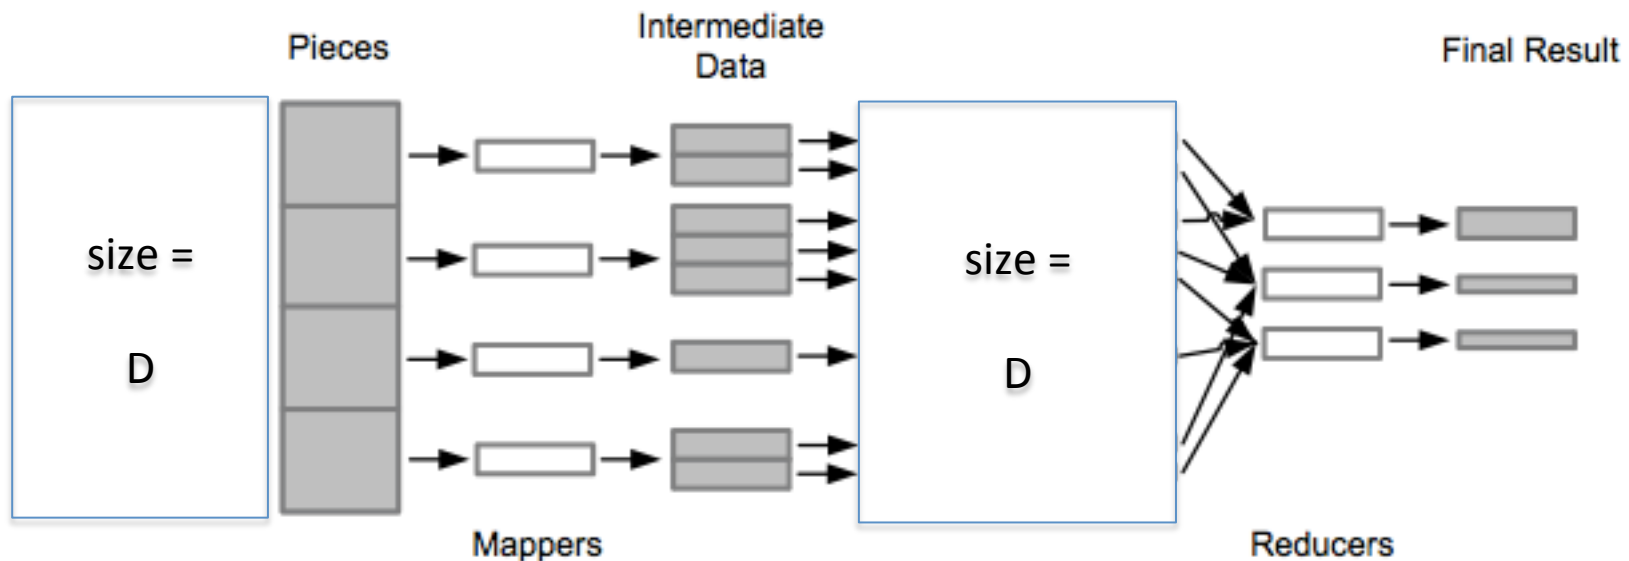


# map, reduce ... also 'combine'


how much data is produced by map?

each word is emitted multiple times!

combiner : sum up word-counts per mapper before emitting



# database join using map-reduce



**SQL:** `SELECT SUM(Sale), City FROM Sales, Cities WHERE Sales.AddrID=Cities.AddrID GROUP BY City`

# real-world example

lots of data ...

paper, author, contents

million such papers, million authors, millions of possible terms ('phrases' occurring in contents)

problems:

top 10 terms for each author; top 10 authors per term...

*'database' person's solution ....*

Q = select id, word, author  
from P where in(w,content)

select count(), word, author  
from Q group by word

| id (paper-id) | content | author |
|---------------|---------|--------|
|               |         |        |
|               |         |        |
|               |         |        |



| id | word | author |
|----|------|--------|
|    |      |        |
|    |      |        |
|    |      |        |




| wc | word | author |
|----|------|--------|
|    |      |        |
|    |      |        |
|    |      |        |

P    million

Q

trillions (million x million)!

# top-k words per author in map-reduce



map:  
emit word, author


reduce:  
reduce-key = word+author  
reduce-function = count

suffers from same problem – trillion combinations!

– map-reduce alone is not enough – *approach needs to change!*



# top-k words per author in map-reduce




map:  
emit author, contents

reduce:  
reduce-key = author  
reduce-function = **F()**

F(): for each author:

scan all inputs and compute word-counts .. insert into w  
sort w, output the top k, delete w and reinitialize to [ ]

# indexing in map-reduce




map:  
produce a partial index  
i.e. emit  $w \rightarrow$  postings-list

reduce:  
reduce-key = word  
merge partial indexes  
i.e. merge postings per word

what about sorting by either document-id, or page-rank etc. ?

# LSH in map-reduce



map:  
emit doc-id, k hash-values

reduce:  
reduce-key = hashes  
emit doc-pairs for each key

will a document-pair be emitted by more than one reducer?

# inside map-reduce

