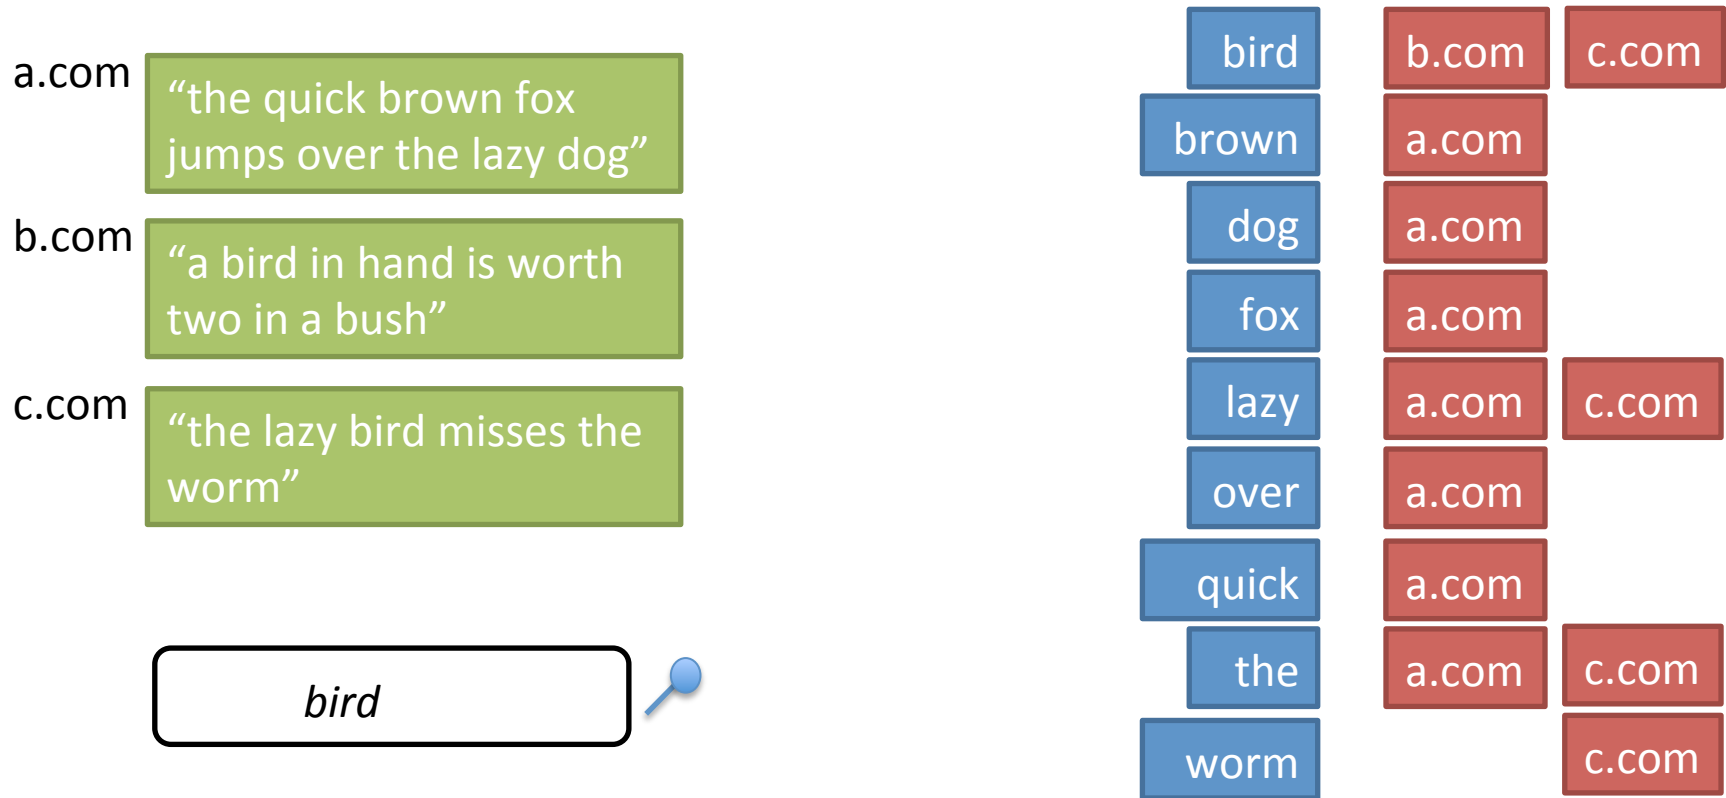


on the web: basic text indexing



looking up a posting takes $O(\log m)$:

keep the term-lists in a sorted structure [hashing - can do better $O(1+m/K)$]

still need to assemble results of a q -term query

$O(rq)$ if $r = \#$ intermediate results in all; what if r is huge ??

how to create a text index

a.com	"the quick brown fox jumps over the lazy dog"
b.com	"a bird in hand is worth two in a bush"
c.com	"the lazy bird misses the worm"

class index:

```
def create(D) :  
    for d in D :  
        for w in d :  
            i = index.lookup(w)  
            if i < 0 :  
                j = index.add(w)  
                index.append(j,d.id)  
            else:  
                index.append(i,d.id)
```

bird	b.com	c.com
brown	a.com	
fox	a.com	
dog	a.com	
lazy	a.com	c.com
over	a.com	
quick	a.com	
the	a.com	c.com
worm		c.com

complexity of index creation

n documents, m words, w words per document

- every word in each document needs to be read, so the complexity is at least $O(n w)$

a.com "the quick brown fox jumps over the lazy dog"

b.com "a bird in hand is worth two in a bush"

c.com "the lazy bird misses the worm"

additionally, as *each* word is read:

- we need to lookup the sorted structure of at most m words to find out if it has already been inserted before; this cost is $O(\log m)$ or $O(1)$ if we use a good hash table
- we must insert the url in the document list for the word (after creating a new entry if needed)

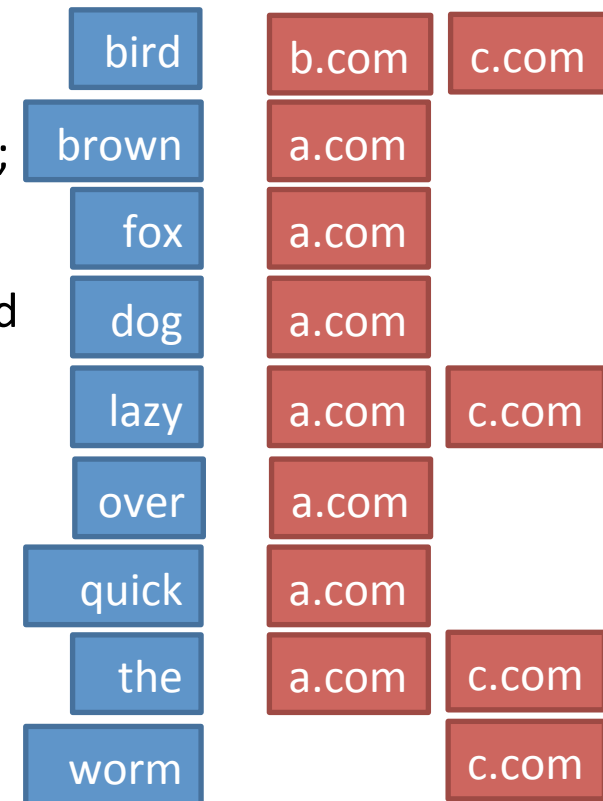
each of these represents but a *constant* cost per word*

**there is an important assumption here – HW...*

therefore the complexity of our procedure is $O(n w \log m)$

(using a balanced binary tree to store words)

or $O(n w)$ (using a hash table to store words)



now that we know what an index is ..

how many web-pages are indexed?

2-5 billion

✓ 30-40 billion

200-300 billion

trillions

search for a common word, such as 'a', or 'in' on Google and see how many results are returned

how to arrange the results of search?

what if the result set is very large?

- e.g. search for `a' in Google
- also – how to assemble results of a *q-term* query
 $O(r\ q)$ if $r = \#$ intermediate results in all;
- search for `Clinton plays India cards':
“Clinton to visit India but Islamabad was not on the cards...”
OR “Clinton Cards acquired, will save hundreds of jobs in India ...”

similarity (from search index) vs importance

- *name the first word that comes to mind ...*
starting with “A”? starting with “G”?
are some words more important than others; just the common words?
- *top 10 documents matching `Clinton plays India cards'*
importance = PageRank + but is there anything deeper?

page rank

imagine a `random surfer`
what is the relative probability of
visiting a *particular* page?


= page-rank of the page

is the number of hyper-links of a
page sufficient to compute its
page-rank?

yes

✓ no

no – because the surfer can re-visit
a page via cycles in the graph
page-rank is a global property



page-rank is computed iteratively,
continuously and in parallel
page-rank is related to the largest
eigenvector of an adjacency matrix

page rank and memory

search results ordered by page-rank have proved 'intuitive' (\Rightarrow \$\$)
does page-rank provide more insight, say into human memory?



"Google and the Mind" Psychological Science, 2007

1. people asked to form
word-word associations

➤ *a semantic network*

2. people asked to form
letter-word associations

Q: could human response
in 2. be predicted from
the semantic net of 1.?



found in top k percentile using algorithm

page-rank did best
*does this mean
anything?*

search vs. memory

is human memory similar to Google's massive 'index'?

yes

✓ no

most of us are poor at remembering facts

"when was Napoleon's defeat at Waterloo?"

we often need context to augment recall

not recognizing a work colleague when seen in a mall ...

memories are linked in time

what one did first thing in the morning ... and thereafter, etc.

an incident from one's first day at school / college / work ...

memories are 'fuzzy' – *can you recall every item in your room?*

can be triggered by very sparse matches – *such as a mere smell*

Google and the mind: co-evolution?

page-rank is intuitive, so the more we rely on it
how does this affect accuracy of page-rank?

page-rank gets better

✓ *page-rank gets worse*

no effect at all

page-rank relies on hyperlinks

why include hyperlinks? easier to just `Google' anything!

so newer pages have fewer hyperlinks: bad for page-rank ☒

we find it hard to remember facts, so we increasingly use Google

if our supposedly associative memories rely on building
associations, which are strengthened when traversed during recall

➤ the more we use Google the less we can remember! ☒

“The Shallows: What the Internet is doing to our Brains”, Nicholas Carr, 2010

Google and the mind: co-evolution? ☒

`mere' indexing is poor at capturing deeper associations between documents, words, and `concepts'

however, as we search and retrieve, we also divulge information on the relative *relevance* of search-results vis-à-vis a query

exploiting such relevance feedback can improve search (augmenting page-rank) ☒

what about us?

exercising recall abilities is not the only time connections are built

we use and create fresh connections when *reasoning*

but reasoning relies on a lot of *facts*

and Google provides these abundantly and easily, encouraging more reasoning, so building more, probably deeper associations! ☒

desktops, email, etc. - `private' search

✓ indexing works

➤ but what about relevance?

- no *links* => cannot directly use page-rank

➤ need to capture and use other associations

- named entities (people, places)
- relevance feedback (by tracking user behavior)

➤ duplicate detection and handling

- multiple versions / formats of the same document

□ is 'search' the only paradigm?

- topic & activity mining, contextual suggestions

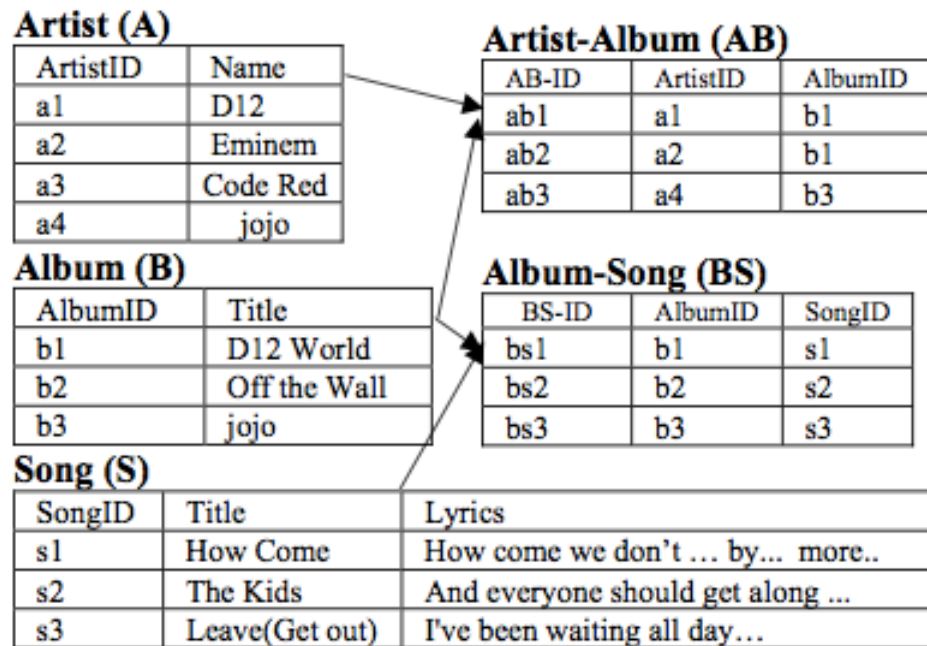
databases & `enterprise search`

all the challenges of `private` search and more:

- context includes the role being played
 - *people play multiple roles*
- taxonomies and classification:
 - manual vs automatic; combinations?
- what about security – role-based access...
- what about `structured` data
 - *SQL is not an answer: text in structured records, linking unstructured documents to structured, `searching` structured records and getting a list of `objects`, i.e. related records*

searching structured data

consider a *LYRICS* database:



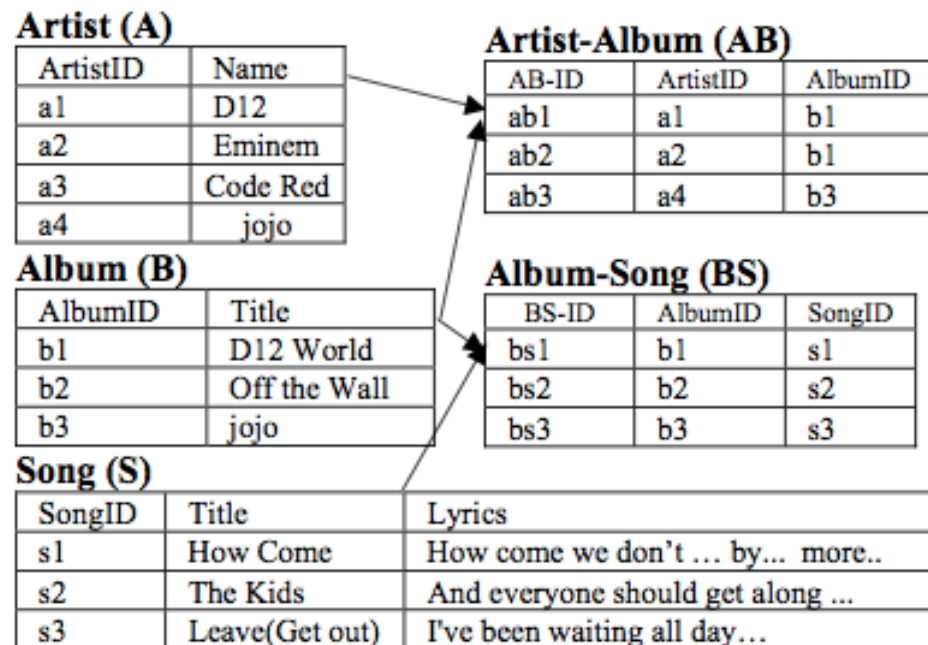
SQL to get albums with “*World*” in the title:

```
Select * from Album B
Where Contains (B.title, 'World' 1) > 0
Order by score(1) desc
```

quiz: searching structured data

how many SQL queries will it take to retrieve

the *names* of each artist and the *lyrics* of every song in an album that has “*World*” in its title



quiz: searching structured data

how many SQL queries will it take to retrieve

the *names* of each artist and the *lyrics* of every song in an album that has “*World*” in its title

Artist (A)

ArtistID	Name
a1	D12
a2	Eminem
a3	Code Red
a4	jojo

Artist-Album (AB)

AB-ID	ArtistID	AlbumID
ab1	a1	b1
ab2	a2	b1
ab3	a4	b3

Album (B)

AlbumID	Title
b1	D12 World
b2	Off the Wall
b3	jojo

Album-Song (BS)

BS-ID	AlbumID	SongID
bs1	b1	s1
bs2	b2	s2
bs3	b3	s3

Song (S)

SongID	Title	Lyrics
s1	How Come	How come we don't ... by... more..
s2	The Kids	And everyone should get along ...
s3	Leave(Get out)	I've been waiting all day...

Query 1: ‘World’ from Album *

Query 2: “lyrics how come by D12”

Query 3: “album by D12 and Eminem”

searching structured data

compare writing SQLs with issuing a ‘search’ query:
“off the world”

- partial matches are missed, e.g. *“World”* , *“off the wall”*
- schema needs to be understood
- many queries, or a complex join are needed

but there is more:

- suppose there were multiple databases, each with a different schema, and with partial, or duplicated data?
 - most important – some unstructured data in documents, other structured in databases: how to search both *together*
- *‘searching’ structured data well remains a research problem*

other kinds of search

index a object (document) by features (words)

assumption is that query is a bag of words, i.e. features

what if the query is an object

e.g. an image (Google Goggles), fingerprint + iris (UID) ...*

is an inverted index the best way to search for objects?

yes

✓ *no*

why? – think about this and discuss!

there is another, very powerful method, called:

Locality Sensitive Hashing**

“compare n pairs of objects in $O(n)$ time”

**Indyk and Motwani '98; Ullman and Rajaraman, Ch 3

*<http://uidai.gov.in/>

locality sensitive hashing (LSH)

basic idea – object x is hashed $h(x)$ so that

if $x = y$ or x close-to y , then $h(x) = h(y)$ with *high probability*,
and conversely

if $x \neq y$ (x far-from y) then $h(x) \neq h(y)$ with *high probability*

constructing the hash functions is tricky ...

combining random functions from a “locally sensitive” *family*

see Ullman and Rajaraman – Chapter 3


example application: biometric matching

e.g. UID, of a billion+ people, *280+ million* enrolled so far ...*

*disclaimer: what UID uses is proprietary, this is merely a motivating example

combining locality-sensitive functions

$$1 - (1 - (pq)^k)^b$$



pq is the probability of a match in one function; even if moderate the LSH expression *amplifies* this match probability while driving the false-match probability to zero as long as it is reasonably smaller

some `big data' applications of LSH

grouping similar tweets without comparing all pairs

near-duplicates / versions of the same root document

finding patterns in time-series (e.g. sensor data)

resolving identities of people from multiple inputs

...

LSH and ‘dimensionality reduction’

intuition

the ‘space’ of objects (prints) is d -dimensional, (e.g. 1000)

2^d , i.e., lots ... of possible objects

LSH *reduces* the dimension to just b hash values (e.g. 1024),


further, random hash functions turn out to be locality sensitive*

so similar objects map to ‘similar’ hash values

- closely related to other kinds of ‘dimensionality-reduction’
- bit tricky to implement, especially in parallel - ...

LSH-based indexing

it might appear that LSH 'groups' similar items
instead it computes the *neighborhood* of each item:
e.g. – represent each object (print) by its b hash-values



approximate recall: associative memory

do we *store* all objects (images, experiences ...)?

“sparse distributed memory” *

pre-dating LSH; also related to high-dimensional spaces, *exploits* vs reduce

consider the space of all 1000-bit vectors; there are lots .. 2^{1000} !

average distance between any two 1000-bit vectors? 500

now – consider a particular vector x chosen at random

half of all other vectors differ by < 500 bits, half by more .. obvious

how many differ from x by less than 450 bits?

binomial distribution with mean 500, $n=1000$, so $\sigma = \sqrt{npq} = \sqrt{250} = 15.8$


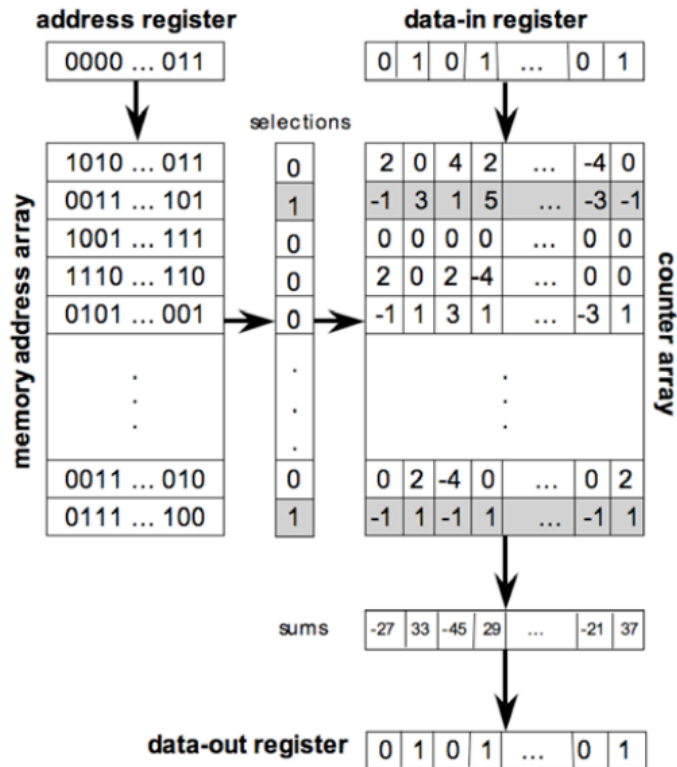
using a normal approximation – only $.0007^{th}$ are less than 450 bits from x

or, most vectors ($.998$, all but $< 2/1000^{ths}$), are within 450 and 550 bits away!

in SDM, concepts are represented by m random vectors:

- ‘nearby’ instances, i.e., even differing in 400 bits, are easily identified
- moreover, SDM shows how to recall by construction – instances accumulate rather than being individually stored

sparse-distributed memory at work



P. J. Denning, American Scientist 77 (July-August 1989)

observed 'documents', 'images' or 'objects' are *not* stored

instead these are *reconstructed* 'from memory'

SDMs can store objects address by *themselves*

SDMs can store *sequences* of objects, addressed by preceding elements

`looking' vs searching

- seeing: recognizing objects and activities
- browsing a bookshelf, flipping pages of a book
- looking at data: time-series, histogram, charts

“*visualizing* a scene”

“getting a *feel* for a document, collection or data”

- ☐ clustering and classification
- ☐ topic discovery, summarization
- ☐ correlation and `interestingness’