

# Open Elective - M.Tech II Semester ICT 586: Big Data Analytics and Technologies

Module-I

## Big Data Overview

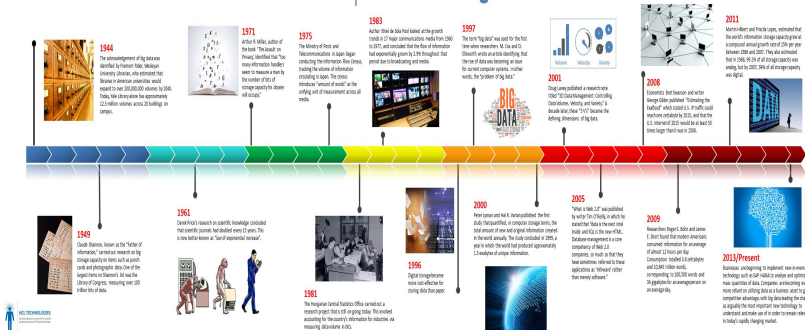
**Girija V. Attigeri**

Dept. of Information & Communication Technology, MIT, Manipal

# Contents

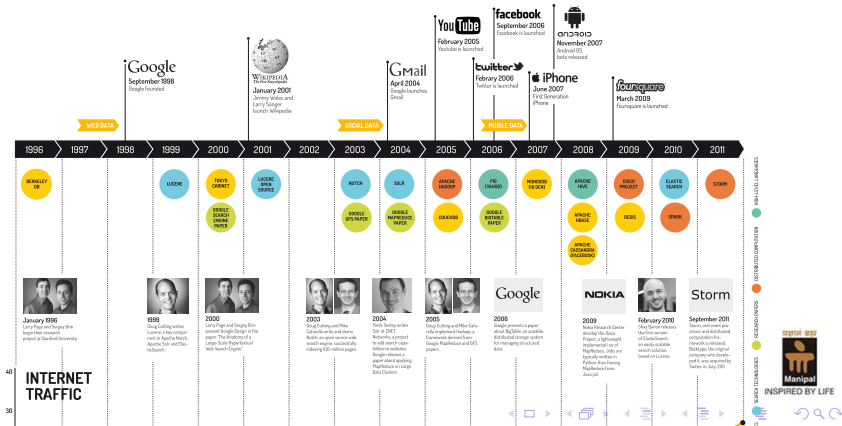
- Introduction to big data.
- Analysis of data at Rest- Hadoop analytics
- Limitations of existing distributing systems - Hadoop Approach
- Hadoop Architecture: Distributed file system: HDFS and GPFS
- MapReduce programming paradigm
- Internals of Hadoop MR engine
- Need for High level language- JAQL and PIG

## Timeline | 70 Years of Big Data



# BIG DATA

## A BRIEF HISTORY




- Air Bus A380
  - 1 billion line of code
  - each engine generate 10 TB every 30 min
  - 640TB per Flight
- Twitter Generate approximately 12 TB of data per day
- New York Stock Exchange 1TB of data everyday

There are huge volumes of data in the world:

- From the beginning of recorded time until 2003, We created 5 billion gigabytes (exabytes) of data.
- In 2011, the same amount was created every two days
- In 2013, the same amount of data is created every 10 minutes.

- How big is the Big Data?
- What is big today maybe not big tomorrow
- Any data that can challenge our current technology in some manner can consider as Big Data
  - Volume
  - Communication
  - Speed of Generating
  - Meaningful Analysis





**639,800 GB of global IP data transferred**

**135** Botnet infections

**6** New Wikipedia articles published

**1,300** New mobile users

**20** New victims of identity theft

**204 million** Emails sent

**47,000** App downloads

**\$83,000** In sales

**277,000** Logins

**6 million** Facebook views

**2+ million** Search queries

**30** Hours of video uploaded

**1.3 million** Video views

**61,141** Hours of music

**20 million** Photo views

**3,000** Photo uploads

**320+** New Twitter accounts

**100,000** New tweets

**And Future Growth is Staggering**

intel

## And Future Growth is Staggering





## Big Data Vectors (3Vs):

Big Data are high-volume, high-velocity, and/or high-variety information assets that require new forms of processing to enable enhanced decision making, insight discovery and process optimization

- Volume: Enterprises are flooded with ever-growing data of all types, easily gathering terabytes even petabytes of information.
- Turingn 12 terabytes of Tweets created each day into improved product sentiment analysis
- Converting 350 billion annual meter readings to better predict power consumption

- Velocity: Sometimes 2 minutes is too late. For time-sensitive processes such as catching fraud, big data must be used as it streams into your enterprise in order to maximize its value.
- Scrutinizing 5 million trade events created each day to identify potential fraud
- Analyzing 500 million daily call detail records in real-time to predict customer churn faster

- Variety: Big data is any type of data - structured and unstructured data such as text, sensor data, audio, video, click streams, log files and more. New insights are found when analyzing these data types together.
- Monitoring 100s of live video feeds from surveillance cameras to target points of interest
- Exploit the 80% data growth in images, video and documents to improve customer satisfaction



- high-volume: Amount of data
- high-velocity: Speed rate in collecting or acquiring or generating or processing of data
- high-variety: Different data type such as audio, video, image data (mostly unstructured data)

Big-Data is similar to Small-data but bigger.. But having data bigger it requires different approaches: Techniques, tools, architecture with an aim to solve new problems Or old problems in a better way



# Analysis of data at Rest Hadoop analytics

The Apache Hadoop software library is a framework that allows for the distributed processing of large data sets across clusters of computers using simple programming models. It is designed to scale up from single servers to thousands of machines, each offering local computation and storage.

## Core components of Hadoop



**hadoop**

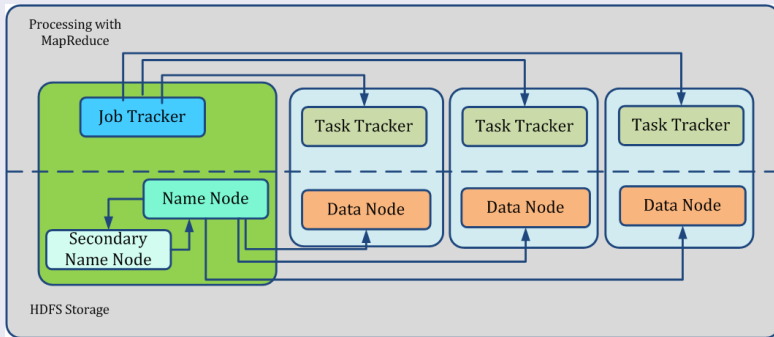
**MapReduce**

Distributed Data Processing

**HDFS**

Distributed Filesystem

## Hadoop Logical Architecture



# Distributed File System

- Distributed File System is a special case of distributed system.
- In the Distributed file system, storage resources and clients are dispersed in the network.
- The main goal of Distributed File System is to provide common view of centralized file system, even though it has a distributed implementation.

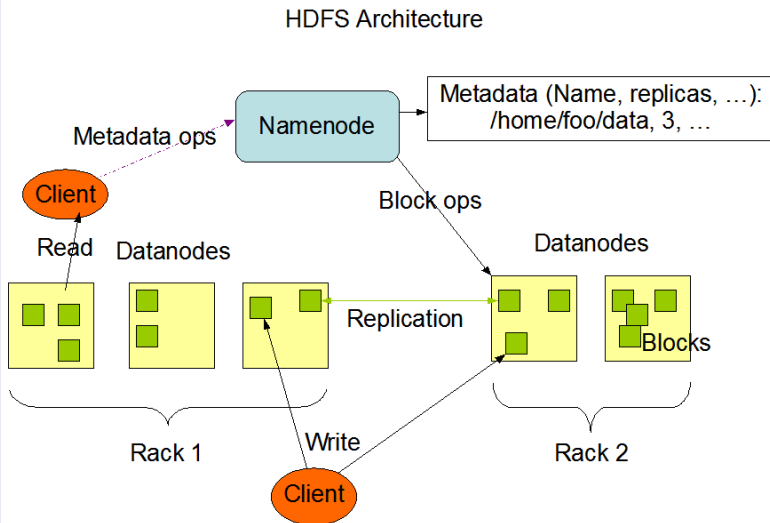
- HDFS is a Java-based file system that provides scalable and reliable data storage, and it was designed to span large clusters of commodity servers.
- HDFS has demonstrated production scalability of up to 200 PB of storage and a single cluster of 4500 servers, supporting close to a billion files and blocks.

- Goal:
  - Batch Processing
  - High aggregate data bandwidth
  - support millions of files in single instance
  - Detection of faults and fast recovery
- Architecture : Master/Slave architecture
- Processes : Stateful
- Communication: TCP/IP
  - Client to name node through configurable TCP port
  - Name Node to Data nodes Data node protocol
  - RPC wraps both Client and Data node protocols
- Naming
- Synchronization: Write Once and read many access
- Caching and Replication:
  - Block size and Replication factor is configurable
  - Rack aware replica policy
- Fault Tolerance

# Assumptions and Goals

- Hardware Failure : detection of faults and quick, automatic recovery from them is a core architectural goal of HDFS
- Streaming Data Access
- Large Data Sets
- Simple Coherency Model
- Moving Computation is Cheaper than Moving Data
- Portability Across Heterogeneous Hardware and Software Platforms

## HDFS Logical Architecture






## Data Replication

### Block Replication

Namenode (Filename, numReplicas, block-ids, ...)  
/users/sameerp/data/part-0, r:2, {1,3}, ...  
/users/sameerp/data/part-1, r:3, {2,4,5}, ...

### Datanodes



## • Replica Placement

- The placement of replicas is critical to HDFS reliability and performance.
- The purpose of a rack-aware replica placement policy is to improve data reliability, availability, and network bandwidth utilization
- When the replication factor is three HDFSs placement policy is
  - put one replica on one node in the local rack
  - another on a node in a different (remote) rack,
  - the last on a different node in the same remote rack
- This policy cuts the inter-rack write traffic which generally improves write performance. The chance of rack failure is far less than that of node failure;

## • Replica Selection

## • Safemode

# The Communication Protocols

- **Robustness :** The primary objective of HDFS is to store data reliably even in the presence of failures. The three common types of failures are NameNode failures, DataNode failures and network partitions.
  - Data Disk Failure, Heartbeats and Re-Replication
  - Cluster Rebalancing
  - Data Integrity
  - Metadata Disk Failure
  - Snapshots

# Data Organization

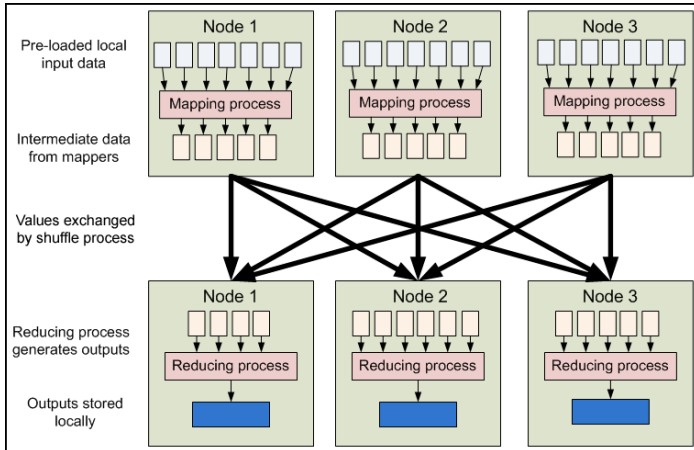
- Data Blocks
- Staging
- Replication Pipelining


## Map Reduce

MapReduce is the heart of Hadoop. It is this programming paradigm that allows for massive scalability across hundreds or thousands of servers in a Hadoop cluster.


The MapReduce concept is fairly simple to understand for those who are familiar with clustered scale-out data processing solutions.

- Users specify the computation in terms of a map and a reduce function,
- the underlying runtime system automatically parallelizes the computation across large-scale clusters of machines, handles machine failures, and schedules inter-machine communication to make efficient use of the network and disks









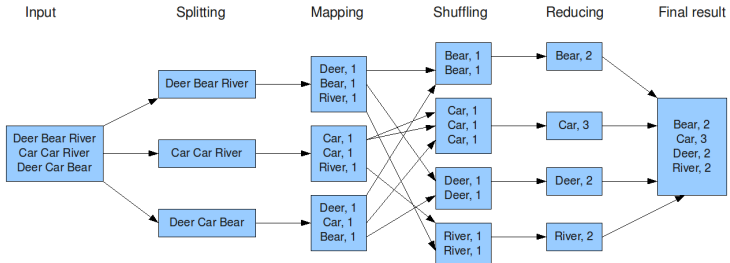
## Mapper

- Mapper maps input key,value pairs to a set of intermediate key,value pairs.
- Maps are the individual tasks that transform input records into intermediate records
- The transformed intermediate records do not need to be of the same type as the input records. A given input pair may map to zero or many output pairs
- The Hadoop MapReduce framework spawns one map task for each InputSplit generated by the InputFormat for the job
- Mapper implementations are passed the JobConf for the job via the JobConfigurable.configure(JobConf)
- Output pairs do not need to be of the same types as input pairs. A given input pair may map to zero or many output pairs. Output pairs are collected with calls to OutputCollector.collect(WritableComparable,Writable).
- Reporter to report progress
- The Mapper outputs are sorted and then partitioned per Reducer. The total number of partitions is the same as the number of reduce tasks for the job.

## Reducer

- reduces a set of intermediate values which share a key to a smaller set of values.
- Reducer has 3 primary phases: shuffle, sort and reduce.
  - Shuffle :Input to the Reducer is the sorted output of the mappers. In this phase the framework fetches the relevant partition of the output of all the mappers, via HTTP.
  - Sort:The framework groups Reducer inputs by keys (since different mappers may have output the same key) in this stage. The shuffle and sort phases occur simultaneously; while map-outputs are being fetched they are merged.
  - Reduce :In this phase the reduce(WritableComparable, Iterator, OutputCollector, Reporter) method is called for each key, (list of values) pair in the grouped inputs.

The overall MapReduce word count process



# Hadoop Ecosystem

Hadoop Ecosystem is a group of numerous self-standing software projects that are built around Hadoop, each addressing a variety of problem spaces and meeting different needs

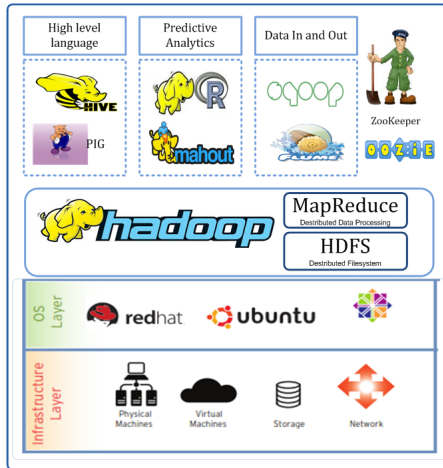


Figure: Hadoop Eco System

# Thank You